



Docker and Kubernetes offline installation in RHEL7

Setup Kubernetes Cluster on Red Hat Enterprise Linux 7 - Offline

Pre-requisites

- **Online Machine:** CentOS 7 machine with internet access
- **Offline Machine:** CentOS 7 machine where Docker & K8s will be installed in offline mode
The idea is to download all the dependencies in the online machine and transfer them to offline machine

Installing Docker

Online Machine

- Login as root
- Configure YUM package manager
`yum update`
`yum install yum-utils device-mapper-persistent-data lvm2`
- Add docker repo
`yum-config-manager --add-repo \`
`https://download.docker.com/linux/centos/docker-ce.repo`
- Install and enable docker
`yum install docker-ce`
`systemctl start docker`
`systemctl enable docker`
- Check docker version
`docker version`
- Create a directory to save all the docker dependencies needed for offline machine
`mkdir ~/docker_dependencies`
`cd ~/docker_dependencies`
`yumdownloader --assumeyes --destdir=$PWD/yum --resolve yum-utils`
`yumdownloader --assumeyes --destdir=$PWD/dm --resolve device-mapper-persistent-data`
`yumdownloader --assumeyes --destdir=$PWD/lvm2 --resolve lvm2`
`yumdownloader --assumeyes --destdir=$PWD/docker-ce --resolve docker-ce`
`yumdownloader --assumeyes --destdir=$PWD/se --resolve container-selinux`
- Copy this folder to offline machine

Offline Machine

- Login as root
- Install all the dependencies downloaded in the previous step
`yum install -y --cacheonly --disablerepo=* docker_dependencies/yum/*.rpm`
`yum install -y --cacheonly --disablerepo=* docker_dependencies/dm/*.rpm`
`yum install -y --cacheonly --disablerepo=* docker_dependencies/lvm2/*.rpm`
`yum install -y --cacheonly --disablerepo=* docker_dependencies/docker-ce/*.rpm`

```
yum install -y --cacheonly --disablerepo=* docker_dependencies/se/*.rpm
```

- Enable and start docker daemon

```
systemctl enable docker
```

```
systemctl start docker
```

```
systemctl status docker
```

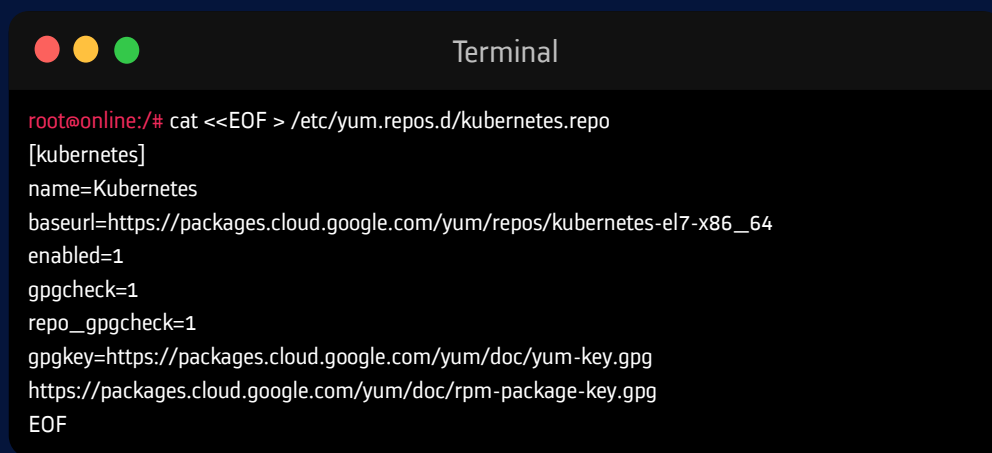
- Check the installation

```
docker version
```

Installing Kubernetes: Setting Up Offline Master Node

Online Machine

- Login as root
- Add Kubernetes yum repository



```

Terminal

root@online:/# cat <<EOF > /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
EOF

```

- Build yum cache

```
yum makecache fast
```
- Download Kubernetes utilities
 - Create a directory to save all the dependencies required by `kubelet`, `kubeadm` & `kubectl`

```
mkdir ~/k8s_dependencies
```

```
cd ~/k8s_dependencies
```

```
yumdownloader --assumeyes --destdir=$PWD --resolve yum-utils kubeadm-1.18.* kubelet-1.18.* kubectl-1.18.* ebttables
```
- Copy this directory to offline machine

Offline Machine: PART 1

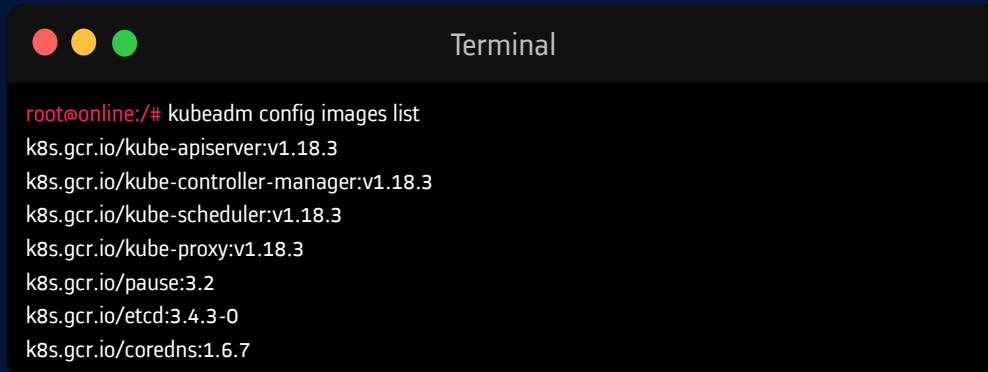
- Login as root
- Install all the dependencies downloaded in the previous step

```
cd k8s_dependencies
```

```
yum install -y --cacheonly --disablerepo=* $PWD/*.rpm
```
- We have successfully installed `kubelet`, `kubeadm` & `kubectl`

- Next step is to find the images required by **kubeadm** to bootstrap the cluster. Run the below command in the terminal to get the list of images

kubeadm config images list



```

root@online:/# kubeadm config images list
k8s.gcr.io/kube-apiserver:v1.18.3
k8s.gcr.io/kube-controller-manager:v1.18.3
k8s.gcr.io/kube-scheduler:v1.18.3
k8s.gcr.io/kube-proxy:v1.18.3
k8s.gcr.io/pause:3.2
k8s.gcr.io/etcd:3.4.3-0
k8s.gcr.io/coredns:1.6.7

```

- We will download all the above images in the online machine and transfer them back to offline machine for bootstrapping the cluster.

Online Machine

- Login as root
- It is assumed that docker is already up and running in this machine
- Pull all docker images required to setup k8s cluster


```

docker pull k8s.gcr.io/kube-apiserver:v1.18.3
docker pull k8s.gcr.io/kube-controller-manager:v1.18.3
docker pull k8s.gcr.io/kube-scheduler:v1.18.3
docker pull k8s.gcr.io/kube-proxy:v1.18.3
docker pull k8s.gcr.io/pause:3.2
docker pull k8s.gcr.io/etcd:3.4.3-0
docker pull k8s.gcr.io/coredns:1.6.7

```
- We will also need the docker image of CNI. For this setup I have used **Flannel** as the CNI
- Find out the image needed by flannel from the below link
<https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>
 It is **quay.io/coreos/flannel:v0.12.0-amd64**
- Download this docker image as well


```

docker pull quay.io/coreos/flannel:v0.12.0-amd64

```
- [Optional] If required, download the images required for setting up Kubernetes dashboard and ingress controllers like nginx, traefik etc., Since offline machines doesn't have access to public internet, feel free to download your docker hub images at this stage. You can do this after the cluster setup as well.
- Save individual images as a TAR archive


```

mkdir k8s_images
cd k8s_images
docker save k8s.gcr.io/kube-apiserver:v1.18.3 > kube-api.tar

```

```

docker save k8s.gcr.io/kube-controller-manager:v1.18.3 > kube-controller.tar
docker save k8s.gcr.io/kube-scheduler:v1.18.3 > kube-sched.tar
docker save k8s.gcr.io/kube-proxy:v1.18.3 > kube-proxy.tar
docker save k8s.gcr.io/pause:3.2 > pause.tar
docker save k8s.gcr.io/etcd:3.4.3-0 > etcd.tar
docker save k8s.gcr.io/coredns:1.6.7 > coredns.tar
docker save quay.io/coreos/flannel:v0.12.0-amd64 > flannel.tar

```

- Transfer this directory with TAR images to offline machine

Offline Machine: PART2

- Login as root
- Unpack the tar images copied from online machine by executing the following command
`docker load < <image name>.tar`
`docker load < kube-api.tar`
 Repeat this for all the images
- Run `docker images` to see all the loaded images
- Disable swap
`swapoff -a`
`sed -e '/swap/s/^/#/g' -i /etc/fstab`
- Switch SELinux to Permissive mode
`setenforce 0`
`sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config`
- Ensure net.bridge.bridge-nf-call-iptables is set to 1 in your sysctl config
`cat <<EOF > /etc/sysctl.d/k8s.conf`
`net.bridge.bridge-nf-call-ip6tables = 1`
`net.bridge.bridge-nf-call-iptables = 1`
`EOF`
- Run `sysctl --system`
- Configure kubectl autocompletion
`echo "source <(kubectl completion bash)" >> ~/.bashrc`
- Allow Kubernetes service ports in Linux firewall
 On master nodes: `firewall-cmd --permanent --add-port={6443,2379,2380,10250,10251,10252}/tcp`
 On worker nodes: `firewall-cmd --permanent --add-port={10250,30000-32767}/tcp`
- Check `kubectl` version
`kubectl version`
- Run the following command to set up the cluster
`kubeadm init --pod-network-cidr=10.244.0.0/16 --apiserver-advertise-address=<IP of Master Node>`

```
kubeadm init --pod-network-cidr=10.244.0.0/16 --apiserver-advertise-address=192.168.56.2
```

Please note 192.168.56.2 is the IP of the interface chosen for the master to listen to incoming requests(api-server). If you have multiple interfaces, choose the IP of the selected interface. Especially required when you have both dynamic and static IP interfaces and you chose to use static IP for avoiding the trouble of initiating the cluster every-time when the master IP changes on dynamic IP interface

Once above command is executed, a token of the form 'kubeadm join 192.168.56.2:6443 --token oqp357.wexmezpiszo5zeeg \

```
--discovery-token-ca-cert-hash
```

```
sha256:67577a18d7abcd1815cf5086e2196dadf5ac39ab3b18fb705694fdafac84580b'
```

will be generated. save this token for later use to join worker nodes in the cluster

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

{These commands will be shown in the terminal once you run kubeadm init...}

- Deploy flannel CNI

```
kubectl apply -f
```

```
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

- Get list of nodes

```
kubectl get nodes
```

Installing Kubernetes: Setting Up Offline Worker Nodes

- Worker nodes will be setup in the same way as the master
- Follow **PART-1** of offline k8s installation which sets up kubelet, kubeadm & kubectl
- Follow **PART-2** of offline k8s installation from disabling swap to checking kubectl version

```

• Disable swap
swapoff -a
sed -e '/swap/s/^/#/g' -i /etc/fstab
• Switch SELinux to Permissive mode
setenforce 0
sed -i 's/^SELINUX=enforcing/SELINUX=permissive/' /etc/selinux/config
• Ensure net.bridge.bridge-nf-call-iptables is set to 1 in your sysctl config
cat <<EOF > /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
EOF
• Run sysctl --system
• Configure kubectl autocompletion
echo "source <(kubectl completion bash)" >> ~/.bashrc
• Allow Kubernetes service ports in Linux firewall
On master nodes: firewall-cmd --permanent --add-port={6443,2379,2380,10250,10251,10252}/tcp
On worker nodes: firewall-cmd --permanent --add-port={10250,30000-32767}/tcp
• Check kubectl version
kubectl version

```

Other Useful Links

Bootstrap Kubernetes cluster using Kubeadm

https://github.com/kunchalavikram1427/Kubernetes_public/blob/master/Bootstrap_K8s_Cluster_Kubeadm.pdf

Learn K8s

https://github.com/kunchalavikram1427/Kubernetes_public/blob/master/Kubernetes_Made_Easy.pdf

